

Week 3: Data Visualization

`{ggplot2}`

Joe Nese

University of Oregon

Fall 2021

Data Visualization with

Week 3

Agenda

Finalize Groups

`{ggplot2}`

- syntax
- continuous data visualizations
- categorical data visualizations
- options
 - color/fill
 - transparency
 - labels
 - facets

Learning Objectives

- Understand the basic syntax requirements for `{ggplot2}`
- Recognize various options for displaying continuous and categorical data
- Familiarity with various `{ggplot2}` options
 - color/fill
 - transparency
 - labels
 - facets

Share!

{datapasta}

- Copy and paste data to and from R
- VERY handy!
- Good for [reprex](#)
 - posting questions on Rstudio Community or stackoverflow

[demo]





The background of the slide is a dark navy blue field filled with numerous small, multi-colored hexagons and dots. The colors include red, yellow, green, blue, orange, grey, and white. These shapes are scattered across the entire background, creating a starry or pixelated effect.

tidyverse

Providing grammar for:

- Graphics
 - `{ggplot2}`
- Data manipulations
 - `{dplyr}`
 - `{tidyr}`
- Expanding area of specialized topics
 - `{lubridate}`
 - `{glue}`
 - `{tidymodels}`
- Many more...

Providing grammar for:

- Graphics
 - `{ggplot2}`
- Data manipulations
 - `{dplyr}`
 - `{tidyr}`
- Expanding area of specialized topics
 - `{lubridate}`
 - `{glue}`
 - `{tidymodels}`
- Many more...

The `{ggplot2}` package

`gg` stands for the "grammar of graphics"

Resources

The `{ggplot2}` package is one of the most popular [R](#) packages, and there are many resources to learn the syntax

- ggplot2 book (email me for digital copy)
- RStudio [cheat sheet](#)
 - Can be helpful, perhaps more so after a little experience
- [R Graphics Cookbook](#)
- [R Graph Gallery](#)

Components

Every ggplot plot has three components

1. data
 - the `data` used to produce the plot
2. aesthetic mappings (`aes`)
 - between variables and visual properties
3. layer(s)
 - usually through the `geom_*()` function to produce geometric shapes to be rendered

`{ggplot}` always takes a data frame (tibble) as the first argument

Basic syntax

```
ggplot(data, aes(x = xvar, y = yvar )) +  
  geom_function()
```

`ggplot()` = the function is `ggplot` and the package is `{ggplot2}`

`data` = the data to be plotted

`aes(x = xvar, y = yvar)` = the `aesthetic` mappings

`geom_function()` = the `geometrics` of the plot; the "function" here represents any of the geom offerings

`+` = note the `+` and NOT the `%>%`

{ggplot2} template

```
ggplot("data", aes("mappings")) +  
  "geom_function"()
```

or


```
"data" %>%  
  ggplot(aes("mappings")) +  
  "geom_function"()
```

Some data for today

penguins from {palmerpenguins}

[run the following]


```
## # A tibble: 6 x 8
##   species island    bill_length_mm bill_depth_mm flipper_length_mm body_mas
##   <fct>   <fct>          <dbl>         <dbl>         <int>      <i
## 1 Adelie  Torgersen          39.1           18.7           181        3
## 2 Adelie  Torgersen          39.5           17.4           186        3
## 3 Adelie  Torgersen          40.3            18           195        3
## 4 Adelie  Torgersen          NA            NA            NA
## 5 Adelie  Torgersen          36.7           19.3           193        3
## 6 Adelie  Torgersen          39.3           20.6           190        3
```



Continuous Data


Setting up a plot

Run the following code. What do you see?



It's ready for you to add some **geometric layers**...what should we add?

How about points?



Adding layers

- In the previous slide, we added a layer of points
- The `geom_point()` layer is a function, complete with it's own arguments


Let's change the color of the points

How would you change the color of the points?

or


How would you find out about changing the color of the points?

Color



Add another layer

Let's add a smoothed line with `geom_smooth()`



You try

You probably got the **message** below when you ran (defaults)

```
### `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

Change the **method** to **"lm"**

Let's do this together

Look at the help page – [?geom_smooth](#)

1. Remove the confidence interval around the line
2. Now change the *SE* band to reflect a 68% confidence interval

color: global vs. conditional

Prior examples changed colors globally


- `geom_point(color = "magenta")`

Use `aes()` to access variables, and color according to a specific variable

- We use variable names within `aes()`

Let's check the data again (`head()`) and the "species" variable (`table()`)

[let's do this together]



color: global vs. conditional

- When we did `geom_point(color = "magenta")` we put quotes around the color
- Why now is "species" not in quotes?
 - color names/hex codes are in quotes **NOT** in the `aes()`
 - variable names are in the `aes()` **NOT** in quotes
 - `aes()` is where you map to your data!

Conditional flow through layers


If we use something like `color = "x"` in the first `aes`thetic, it will carry on through all additional layers

- These two codes are the same:
- But these two are not...why? [\[run to find out\]](#)

Be mindful with aes()

Using `aes()` when you **don't** need it

What is happening here?



Be mindful with aes()

Not using `aes()` when you need it

What is happening here?

```
### Error in layer(data = data, mapping = mapping, stat = stat, geom = GeomPoin
```

Kind of helpful message here.

Themes

Let's talk themes

- The default is `theme_gray`
 - I don't like it
- But there are a lot of build-in alternative in `{ggplot2}`
 - `theme_minimal` is my favorite
- Check out the `{ggthemes}` package for a lot of alternatives
 - These days I nearly always use the `colorblind` theme for discrete values in my plots
- Check out the `{ggthemeassist}` add-in

More themes


- The `{hrbrthemes}` are nice
- Consider [building your own theme!](#)
- Or Google around
- Set the theme globally
 - One of the first lines in your .Rmd file could be:


```
theme_set(theme_minimal())
```


- I put this in the `setup` chunk


Get a little fancy

- You can use `geom_point()` for more than one layer
- You can also use a different data source on a layer
- Use these two properties to highlight points
 - How about penguins from Torgersen Island?



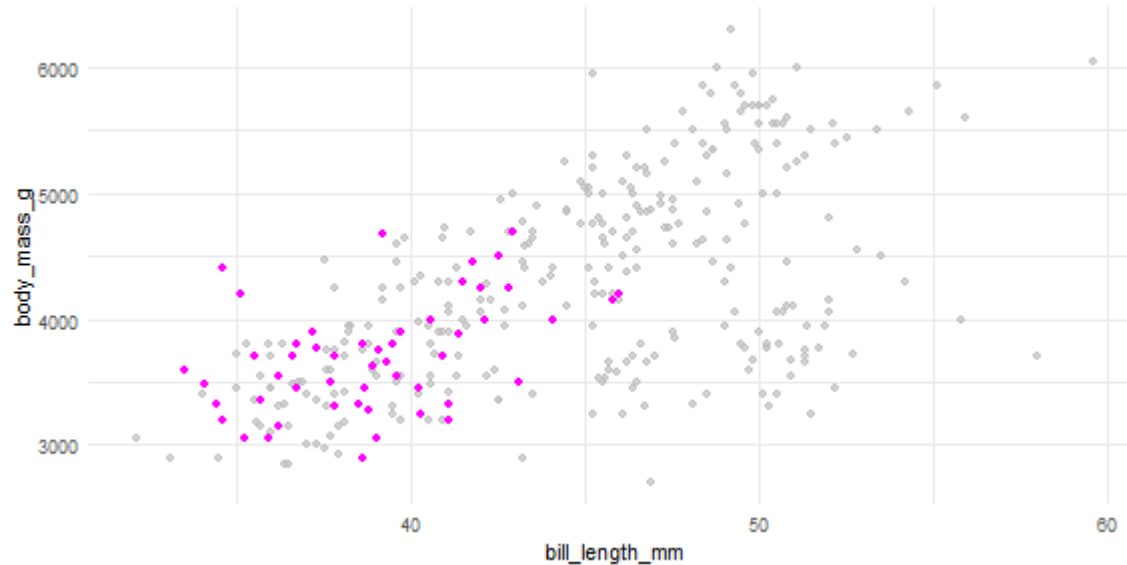






Another option

`{gghighlight}` varying flexibility



Line plots

- When should you use line plots instead of smooths?
 - usually when time is involved
- What are some good candidate data for line plots?
 - observed versus model-implied (estimated)

geom_line()


Classic time series example

economics data from `{ggplot2}`

```
## # A tibble: 574 x 6
##   date      pce    pop psavert uempmed unemploy
##   <date>    <dbl> <dbl>   <dbl>   <dbl>   <dbl>
## 1 1967-07-01  507. 198712   12.6     4.5    2944
## 2 1967-08-01  510. 198911   12.6     4.7    2945
## 3 1967-09-01  516. 199113   11.9     4.6    2958
## 4 1967-10-01  512. 199311   12.9     4.9    3143
## 5 1967-11-01  517. 199498   12.8     4.7    3066
## 6 1967-12-01  525. 199657   11.8     4.8    3018
## 7 1968-01-01  531. 199808   11.7     5.1    2878
## 8 1968-02-01  534. 199920   12.3     4.5    3001
## 9 1968-03-01  544. 200056   11.7     4.1    2877
## 10 1968-04-01  544 200208   12.3     4.6    2709
## # ... with 564 more rows
```

Let's try it

How do you think we'd fit a line plot to these data, showing unemployment ("unemploy") over time?



Layers


What happens when we layer `geom_line` and `geom_point`?

try it!


Not the best instance of this It would work better on a plot with fewer time points, but you get the idea

Labels


Axis Labels



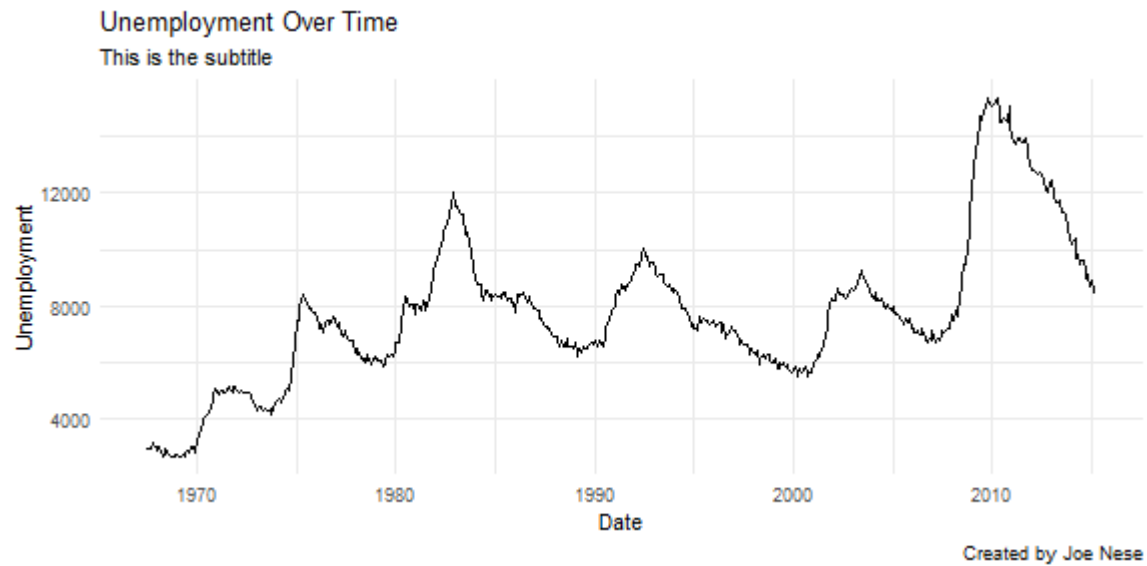
Title



Subtitle



Caption




Tag

(A)


Unemployment Over Time

This is the subtitle



Created by Joe Nese

Legend (one way)




Facets

Faceting

- One of the most powerful features of `{ggplot}`
- Produce n plots by a specific variable
- `facet_wrap()`
 - wrap a sequence of panels into two dimensions
 - based on variables(s)

Faceting

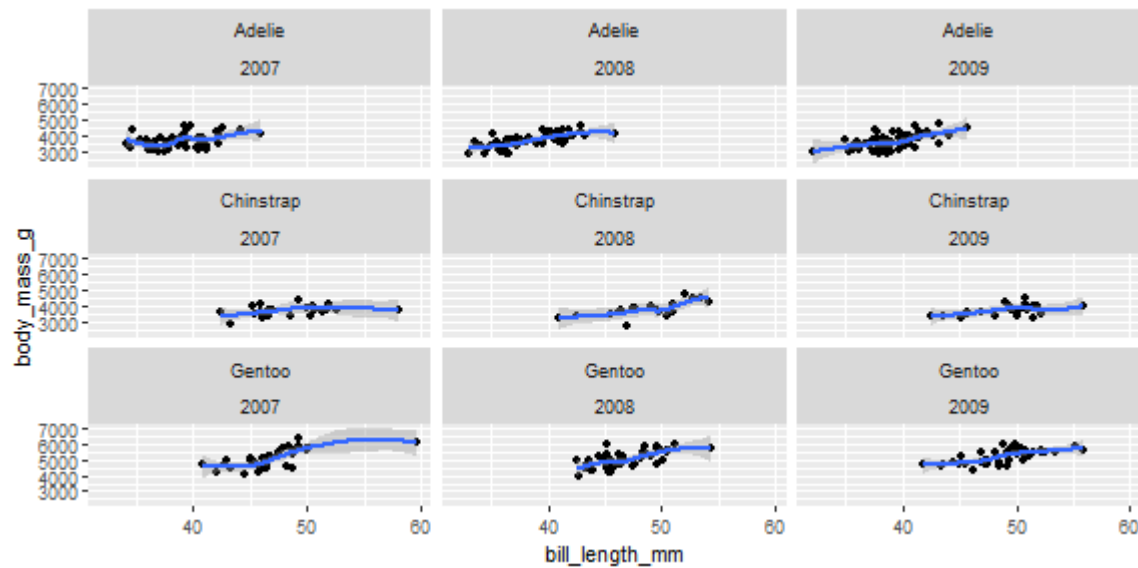


Careful about the ~

```
## Error in validate_facets(x): object 'species' not found
```


Faceting

two variables (like a matrix)



Alternative specification (`vars()`)


Heatmaps


A heatmap is a literal way of visualizing a table of numbers, where you substitute the numbers with colored cells. -- Nathan Yau

- Useful for finding highs and lows - and sometimes patterns
- They don't always work well

Example with correlations

```
## # A tibble: 6 x 3
##   row   col   cor
##   <chr> <chr> <dbl>
## 1 mpg   mpg     1
## 2 mpg   cyl  -0.852
## 3 mpg   disp -0.848
## 4 mpg   hp   -0.776
## 5 mpg   drat  0.681
## 6 mpg   wt   -0.868
```





Categorical Data

Data

`{fivethirtyeight}` package

- `college_grad_students`

```
## # A tibble: 173 x 22
```

```
##   major_code major major_category grad_total grad_sample_size grad_employe
```

```
##           <int> <chr> <chr>           <int>           <int>           <int>
```

```
## 1      5601 Cons~ Industrial Ar~         9173             200             709
```

```
## 2      6004 Comm~ Arts                    53864             882            4049
```

```
## 3      6211 Hosp~ Business                24417             437            1836
```

```
## 4      2201 Cosm~ Industrial Ar~         5411              72             359
```

```
## 5      2001 Comm~ Computers & M~         9109             171             751
```

```
## 6      3201 Cour~ Law & Public ~         1542              22             100
```

```
## 7      6206 Mark~ Business              190996            3738           15157
```

```
## 8      1101 Agri~ Agriculture &~        17488             386            1310
```

```
## 9      2101 Comp~ Computers & M~         5611              98             471
```

```
## 10     1904 Adve~ Communication~        33928             688            2851
```

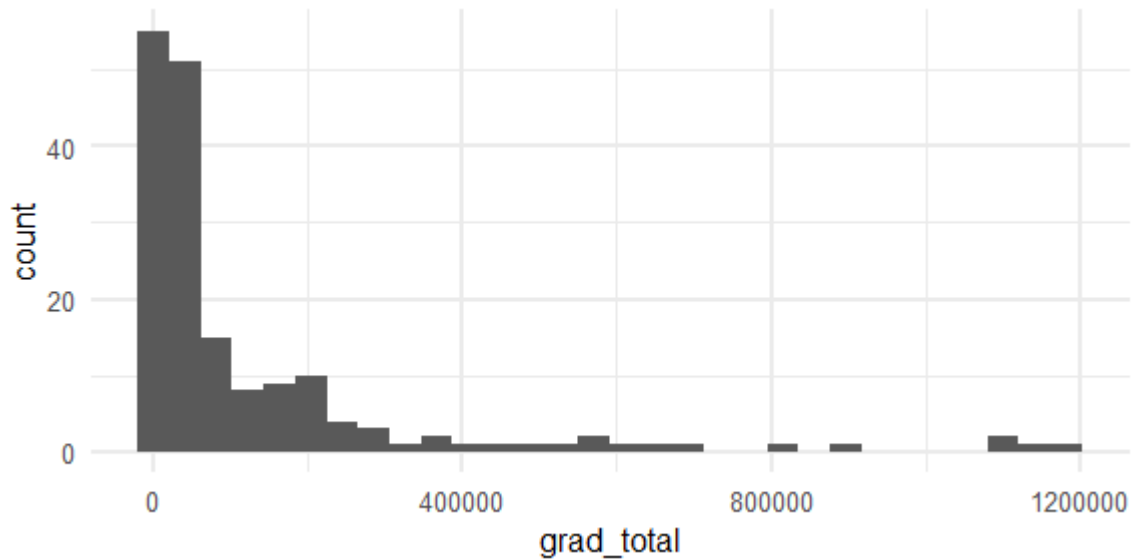
```
## # ... with 163 more rows, and 11 more variables: grad_p75th <int>, nongrad_
```

```
## #   nongrad_employed_fulltime_yearround <int>, nongrad_unemployed <int>, no
```

```
## #   nongrad_p75th <dbl>, grad_share <dbl>, grad_premium <dbl>
```



Histogram

Histogram of "grad_total"



Transparency - alpha

Add some transparency – perhaps this looks nicer



color vs. fill

In general


- `color` defines the color a geom is *outlined*
- `fill` defines the color a geom is *filled*

For example:


- `geom_point()` default has only has a color and **NO** fill because they're just points
- Point shapes 21-24 include both a color and a fill

How would we change the color of this plot?

How would we change the color of this plot?




How would we change the color fill of this plot?



Color by variable

What if we wanted different colors by a variable

- *major_category*



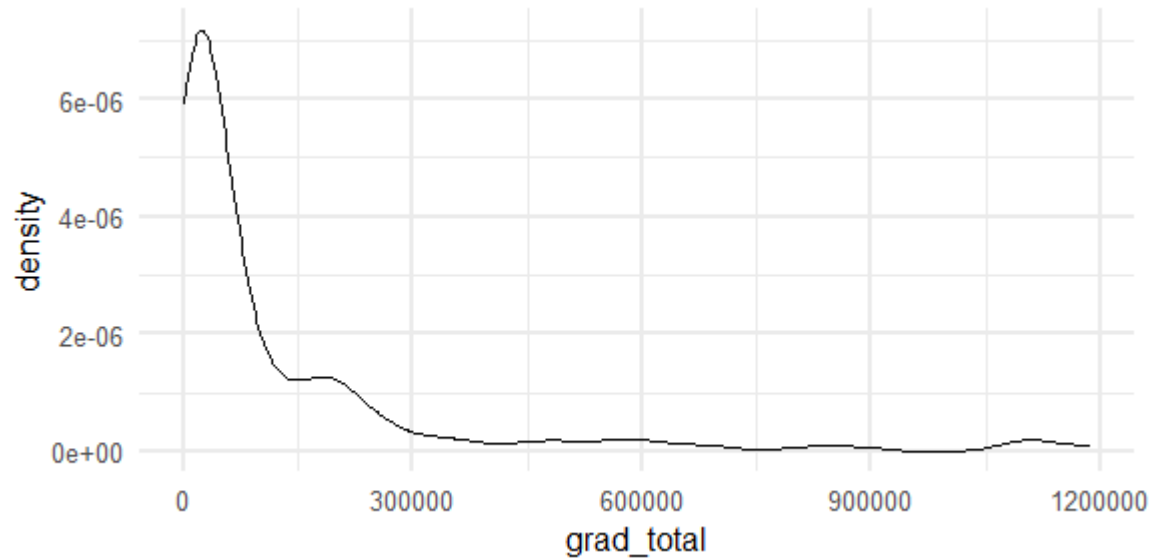
Density plot

Alternative representation of distribution

- Think of it as a smoothed histogram (uses kernel smoothing)
- The depiction of the distribution is **NOT** determined by the number of the bins you use, as are histograms


Density plot

`geom_density()`



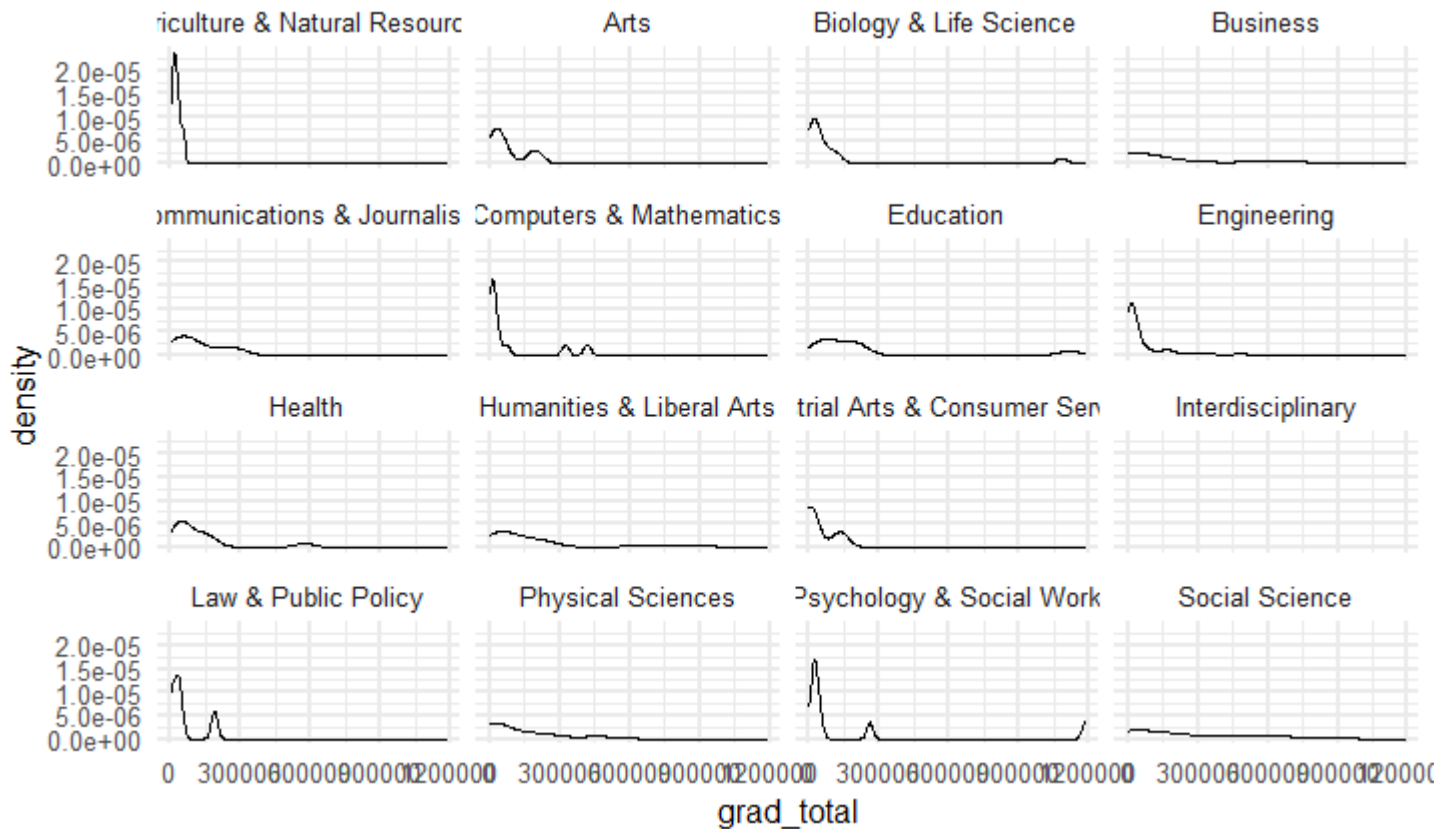
Density plot

Now let's fill by *major_category*




Not a good plot, just an example

Possible alternative? `facet_wrap`




Even better


density ridges




Density Ridges



fill




Add transparency for clarity



Viridis

- easier to read by those with colorblindness
- prints well in gray scale



Same fill function, different option


Candy rankings

{fivethirtyeight}


```
## # A tibble: 197 x 5
```

```
##      competitorname sugarpercent pricepercent winpercent type
##      <chr>          <dbl>          <dbl>          <dbl> <chr>
##  1 100 Grand        0.732          0.860          67.0 chocolate
##  2 100 Grand        0.732          0.860          67.0 caramel
##  3 100 Grand        0.732          0.860          67.0 crispedricewafer
##  4 100 Grand        0.732          0.860          67.0 bar
##  5 3 Musketeers      0.604          0.511          67.6 chocolate
##  6 3 Musketeers      0.604          0.511          67.6 nougat
##  7 3 Musketeers      0.604          0.511          67.6 bar
##  8 Air Heads        0.906          0.511          52.3 fruity
##  9 Almond Joy        0.465          0.767          50.3 chocolate
## 10 Almond Joy        0.465          0.767          50.3 peanutyalmondy
## # ... with 187 more rows
```

Boxplot



Violin plots



Bar Charts

```
## # A tibble: 6 x 11
```

```
## manufacturer model displ year cyl trans      drv      cty      hwy fl      c
##      <chr>      <chr> <dbl> <int> <int> <chr>      <chr> <int> <int> <chr> <
## 1 audi          a4      1.8  1999     4 auto(l5)    f       18     29 p      c
## 2 audi          a4      1.8  1999     4 manual(m5)  f       21     29 p      c
## 3 audi          a4      2    2008     4 manual(m6)  f       20     31 p      c
## 4 audi          a4      2    2008     4 auto(av)    f       21     30 p      c
## 5 audi          a4      2.8  1999     6 auto(l5)    f       16     26 p      c
## 6 audi          a4      2.8  1999     6 manual(m5)  f       18     26 p      c
```

Bar Charts

```
## # A tibble: 6 x 2
## # Groups:   class [6]
##   class          n
##   <chr>        <int>
## 1 2seater         5
## 2 compact       47
## 3 midsize       41
## 4 minivan       11
## 5 pickup        33
## 6 subcompact    35
```

Bar/Col Charts

`geom_bar()`


- expects x **OR** y
- counts rows
- if you want to count the number of cases at each x or y position
- makes the height of the bar proportional to the number of cases in each group
- uses `stat_count()` by default

`geom_col()`

- expects x **AND** y
- expects numbers in your data
- if you want the heights of the bars to represent values in the data
- leaves the data as is
- uses `stat_identity()` by default


geom_bar()

mpg data




geom_col()


summarized_mpg data




geom_bar() default




geom_bar(stat = "identity")




geom_bar() default




geom_col() default



geom_bar() default




geom_bar() uh-oh



What happened?

Let's look at our data again



```
## # A tibble: 7 x 2
## # Groups:   class [7]
##   class      n
##   <chr>    <int>
## 1 2seater      5
## 2 compact    47
## 3 midsize    41
## 4 minivan    11
## 5 pickup     33
## 6 subcompact 35
## 7 suv       62
```



Change the order of the x-axis

- We'll talk about automated methods more later
- Basically - define x-axis variable as a factor with levels in the order you want

```
## # A tibble: 7 x 2
##   class      mean_hwy
##   <chr>      <dbl>
## 1 2seater    24.8
## 2 compact   28.3
## 3 midsize   27.3
## 4 minivan   22.4
## 5 pickup    16.9
## 6 subcompact 28.1
## 7 suv       18.1
```




More bar plot options

Data


```
## # A tibble: 16 x 3
## # Groups:   k_type [2]
##   k_type      ethnic      t1r_mean
##   <chr>      <chr>      <dbl>
## 1 full-day AMERICAN INDIAN OR ALASKA NATIVE      17.5
## 2 full-day ASIAN                               31.8
## 3 full-day BLACK OR AFRICAN AMERICAN, NON-HISPANIC 21.1
## 4 full-day HISPANIC, RACE NOT SPECIFIED            20.3
## 5 full-day HISPANIC, RACE SPECIFIED                22.3
## 6 full-day MORE THAN ONE RACE, NON HISPANIC        20.8
## 7 full-day NATIVE HAWAIIAN, OTHER PACIFIC ISLANDER 23.2
## 8 full-day WHITE, NON-HISPANIC                    25.5
## 9 half-day AMERICAN INDIAN OR ALASKA NATIVE        30.0
## 10 half-day ASIAN                                25.1
## 11 half-day BLACK OR AFRICAN AMERICAN, NON-HISPANIC 18.6
## 12 half-day HISPANIC, RACE NOT SPECIFIED            18.7
## 13 half-day HISPANIC, RACE SPECIFIED                20.0
## 14 half-day MORE THAN ONE RACE, NON HISPANIC        20.3
## 15 half-day NATIVE HAWAIIAN, OTHER PACIFIC ISLANDER 19.7
## 16 half-day WHITE, NON-HISPANIC                    23.7
```


Stacked bar plot

Look for effects in "ethnicity" by "k_type" (full/half day K)




Grouped bar plot




Rotating Labels

I have to look [this](#) up every time




Flip the coordinates

`coord_flip()`




Alternatively

`facet_wrap()`



Sometimes some redundancy works well



geom_*() Review

- `geom_point()`
- `geom_smooth()`
- `geom_line()`
- `geom_tile()`
- `geom_histogram()`
- `geom_density()`
- `ggribes::geom_density_ridges()`
- `geom_boxplot()`
- `geom_violin()`
- `geom_bar()`
- `geom_col()`

Challenge

- Start a new **R** project
- Create a new script, save it as "practice-plots.R"
- Load the `{tidyverse}`
- Print the **msleep** dataset to see its structure (it's from `{ggplot2}`)

For each of the following, produce a separate plot

1. Plot the relation between "sleep_total" and "brainwt" (with "brainwt" as the DV) - scatter plot
2. Overlay a smooth on the previous plot
3. Color the points by "vore", but fit a single smooth
4. Fit separate smooths by "vore", but with all points being gray
5. Omit the standard error of the smooths
6. Use `ylim()` as an additional layer to restrict the y-axis to range from 0 to 5

Next time

Before next class

- Reading
 - [R4DS 5](#)
- Supplemental Learning
 - [RStudio Primer: Working with Tibbles](#)
 - [Rbootcamp: Ch 3](#)
 - [Codecademy: Aggregates in R](#)
- Homework
 - **Homework 2**
 - **Homework 3**

