

# Mutating Joins

**Joe Nese**

University of Oregon

Fall 2021



# Share

`{tidylog}`

- Provides feedback about `{dplyr}` and `{tidyr}` operations

[\[view link\]](#)

# Homework(s) Review

# Mutating Joins

## Week 5

Credit Daniel Anderson for slides

# Agenda

- Quick Share
- `bind_rows()`
- `*_join()`

## Overall Purpose

- Understand and be able to identify keys
- Understand different types of mutating joins
  - `left_join`, `right_join`
  - one-to-one, one-to-many
- Understand some ways joins fail

# A bit about joins

- Also data "merge"
- Today we'll talk about **mutating** joins
- Mutating joins add columns to a dataset
- Mutating joins are the most common, but **filtering** joins can be very powerful

## What if I want to add rows?

- Not technically a join (no key involved)

# Binding rows

```
g3 <- tibble(sid = 1:3,  
             grade = rep(3, 3),  
             score = as.integer(rnorm(3, 200,10)))  
g4 <- tibble(sid = 9:11,  
             grade = rep(4, 3),  
             score = as.integer(rnorm(3, 200,10)))
```

g3

```
## # A tibble: 3 x 3  
##   sid grade score  
##   <int> <dbl> <int>  
## 1     1     3   184  
## 2     2     3   203  
## 3     3     3   212
```

g4

```
## # A tibble: 3 x 3  
##   sid grade score  
##   <int> <dbl> <int>  
## 1     9     4   175  
## 2    10     4   204  
## 3    11     4   193
```



# bind\_rows()

- In examples like the previous data sets, we just want to combine the data by **stacking the rows**
- Data have same (or approximately same) columns
- We can do so with `bind_rows()`

```
bind_rows(g3, g4)
```

```
## # A tibble: 6 x 3
##   sid grade score
##   <int> <dbl> <int>
## 1     1     3   184
## 2     2     3   203
## 3     3     3   212
## 4     9     4   175
## 5    10     4   204
## 6    11     4   193
```

# dplyr::bind\_rows()

- an efficient way to bind many data frames into one, by stacking rows
  - can bind multiple datasets

```
one <- mtcars[1:4, ]  
two <- mtcars[6:10, ]  
three <- mtcars[12:14, ]  
bind_rows(one, two, three)
```

- like joining (merging) data frames that have the same columns
- columns don't need to match when row-binding

# Optional `.id` argument

- What if we knew the grade, but didn't have a variable in each dataset already?
- Use `.id` to add an index for each dataset

```
bind_rows(select(g3, -grade), select(g4, -grade), .id = "dataset")
```

```
## # A tibble: 6 x 3
##   dataset    sid score
##   <chr>    <int> <int>
## 1 1      1     184
## 2 1      2     203
## 3 1      3     212
## 4 2      9     175
## 5 2     10     204
## 6 2     11     193
```

# Recode .id column

```
bind_rows(select(g3, -grade), select(g4, -grade), .id = "dataset") %>%  
  mutate(grade = ifelse(dataset == 1, 3, 4))
```

```
## # A tibble: 6 x 4  
##   dataset  sid score grade  
##   <chr>   <int> <int> <dbl>  
## 1 1      1    184     3  
## 2 1      2    203     3  
## 3 1      3    212     3  
## 4 2      9    175     4  
## 5 2     10    204     4  
## 6 2     11    193     4
```

# Even better

```
bind_rows(select(g3, -grade), select(g4, -grade), .id = "grade") %>%  
  mutate(grade = ifelse(grade == 1, 3, 4))
```

```
## # A tibble: 6 x 3  
##   grade  sid score  
##   <dbl> <int> <int>  
## 1     3     1  184  
## 2     3     2  203  
## 3     3     3  212  
## 4     4     9  175  
## 5     4    10  204  
## 6     4    11  193
```

# What if columns don't match exactly?

Pads with NA

```
bind_rows(g3, g4[, -2], .id = "dataset")
```

```
## # A tibble: 6 x 4
##   dataset  sid grade score
##   <chr>   <int> <dbl> <int>
## 1 1      1      3    184
## 2 1      2      3    203
## 3 1      3      3    212
## 4 2      9     NA    175
## 5 2     10     NA    204
## 6 2     11     NA    193
```

# You can also `bind_cols()`

```
read <- tibble(sid = 1:3,  
               read = as.integer(rnorm(3, 200, 10)))  
math <- tibble(math = as.integer(rnorm(3, 200, 10)))
```

read

math

```
## # A tibble: 3 x 2  
##   sid  read  
##   <int> <int>  
## 1     1   202  
## 2     2   206  
## 3     3   190
```

```
## # A tibble: 3 x 1  
##   math  
##   <int>  
## 1   202  
## 2   202  
## 3   204
```

# bind\_cols()

```
bind_cols(read, math)
```

```
## # A tibble: 3 x 3  
##   sid read math  
##   <int> <int> <int>  
## 1     1    202    202  
## 2     2    206    202  
## 3     3    190    204
```



# Joins

(not to be confused with row binding)

# Keys

- Uniquely identify rows in a dataset
- Variable(s) in common between two datasets to be joined
- A key can be more than one variable

## Types of keys

- Small distinction that you probably won't have to worry about much, but is worth mentioning:
  - **Primary keys**: uniquely identify observations in *their* dataset
  - **Foreign keys**: uniquely identify observations in *other* datasets

# What's the primary key here?

First, let's break down the code:

```
library(rio)
library(here)
ecls <- import(here("data", "ecls-k_samp.sav"),
               setclass = "tbl_df") %>%
  characterize()
```

```
## # A tibble: 984 x 33
```

	child_id	teacher_id	school_id	k_type	school_type	sex	ethnic	famty
	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
##	1	0842021C	0842T02	0842	full-day	public	male	BLACK 0~
##	2	0905002C	0905T01	0905	full-day	private	male	ASIAN
##	3	0150012C	0150T01	0150	full-day	private	female	BLACK 0~
##	4	0556009C	0556T01	0556	full-day	private	female	HISPANI~
##	5	0089013C	0089T04	0089	full-day	public	male	WHITE, ~
##	6	1217001C	1217T13	1217	half-day	public	female	NATIVE ~
##	7	1092008C	1092T01	1092	half-day	public	female	HISPANI~
##	8	0083007C	0083T16	0083	full-day	public	male	WHITE, ~
##	9	1091005C	1091T02	1091	half-day	private	male	WHITE, ~
##	10	2006006C	2006T01	2006	full-day	private	male	WHITE, ~
##	#	... with 974 more rows, and 21 more variables: T1RSCALE <dbl>, T1MSCALE <						

# Let's verify the key

```
ecds %>%  
  count(child_id)
```

```
## # A tibble: 984 x 2  
##   child_id      n  
##   <chr>    <int>  
## 1 0001010C      1  
## 2 0002010C      1  
## 3 0009005C      1  
## 4 0009014C      1  
## 5 0009026C      1  
## 6 0013003C      1  
## 7 0016004C      1  
## 8 0016009C      1  
## 9 0022005C      1  
## 10 0022014C      1  
## # ... with 974 more rows
```

# Let's verify the key

```
ecds %>%  
  count(child_id) %>%  
  arrange(desc(n)) %>%  
  slice(1:3)
```

```
## # A tibble: 3 x 2  
##   child_id      n  
##   <chr>    <int>  
## 1 0001010C      1  
## 2 0002010C      1  
## 3 0009005C      1
```

OR

```
ecds %>%  
  count(child_id) %>%  
  filter(n > 1)
```

```
## # A tibble: 0 x 2  
## # ... with 2 variables: child_id <chr>, n <int>
```

# What about the key here?

```
income_ineq <- read_csv(here("data", "incomeInequality_tidy.csv"))  
head(income_ineq, n = 15)
```

```
## # A tibble: 15 x 6
```

	Year	Number.thousands	realGDPperCap	PopulationK	percentile	income
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	1947	37237	14117.	144126	20	14243
## 2	1947	37237	14117.	144126	40	22984
## 3	1947	37237	14117.	144126	60	31166
## 4	1947	37237	14117.	144126	80	44223
## 5	1947	37237	14117.	144126	50	26764.
## 6	1947	37237	14117.	144126	90	41477
## 7	1947	37237	14117.	144126	95	54172
## 8	1947	37237	14117.	144126	99	134415
## 9	1947	37237	14117.	144126	99.5	203001
## 10	1947	37237	14117.	144126	99.9	479022
## 11	1947	37237	14117.	144126	100.	1584506
## 12	1948	38624	14452.	146631	20	13779
## 13	1948	38624	14452.	146631	40	22655
## 14	1948	38624	14452.	146631	60	30248
## 15	1948	38624	14452.	146631	80	42196

```
income_ineq %>%  
  count(Year, percentile) %>%  
  filter(n > 1)
```

```
## # A tibble: 0 x 3
```

```
## # ... with 3 variables: Year <dbl>, percentile <dbl>, n <int>
```

# Sometimes there is no key

These tables have an *implicit* id - the row numbers. For example:

```
install.packages("nycflights13")  
library(nycflights13)
```

```
head(flights)
```

```
## # A tibble: 6 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_ti  
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <in  
## 1  2013     1     1     517             515           2     830           8  
## 2  2013     1     1     533             529           4     850           8  
## 3  2013     1     1     542             540           2     923           8  
## 4  2013     1     1     544             545          -1    1004          10  
## 5  2013     1     1     554             600          -6     812           8  
## 6  2013     1     1     554             558          -4     740           7  
## # ... with 8 more variables: tailnum <chr>, origin <chr>, dest <chr>, air_t  
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```



```
flights %>%  
  count(year, month, day, flight, tailnum) %>%  
  filter(n > 1)
```

```
## # A tibble: 11 x 6
```

	year	month	day	flight	tailnum	n
	<int>	<int>	<int>	<int>	<chr>	<int>
## 1	2013	2	9	303	<NA>	2
## 2	2013	2	9	655	<NA>	2
## 3	2013	2	9	1623	<NA>	2
## 4	2013	6	8	2269	N487WN	2
## 5	2013	6	15	2269	N230WN	2
## 6	2013	6	22	2269	N440LV	2
## 7	2013	6	29	2269	N707SA	2
## 8	2013	7	6	2269	N259WN	2
## 9	2013	8	3	2269	N446WN	2
## 10	2013	8	10	2269	N478WN	2
## 11	2013	12	15	398	<NA>	2

# Create a key

If there is no key, it's often helpful to add one

These are called *surrogate* keys

```
flights2 <- flights %>%  
  rowid_to_column()  
  
flights2 %>%  
  select(1:3, ncol(flights))
```

```
## # A tibble: 336,776 x 4  
##   rowid  year month minute  
##   <int> <int> <int>   <dbl>  
## 1     1  2013     1     15  
## 2     2  2013     1     29  
## 3     3  2013     1     40  
## 4     4  2013     1     45  
## 5     5  2013     1      0  
## 6     6  2013     1     58  
## 7     7  2013     1      0  
## 8     8  2013     1      0  
## 9     9  2013     1      0
```

# Mutating joins

# Mutating `*_joins()`

- In `{tidyverse}`, we use `mutate()` to create new variables within a dataset
- A mutating join works similarly, in that we're adding new variables to the existing dataset through a join
- **Join**: Two tables of data joined by a common key

# Four types of joins

- `left_join`: keep all the data in the **left** dataset, drop any non-matching cases from the right dataset
- `right_join`: keep all the data in the **right** dataset, drop any non-matching cases from the left dataset
- `inner_join`: keep **only** data that matches in **both** datasets
- `full_join`: keep **all** the data in **both** datasets (also sometimes referred to as an *outer join*)

# Four types of joins

## Mutating joins

- `left_join`: keep all the data in the **left** dataset, drop any non-matching cases from the right dataset
- `right_join`: keep all the data in the **right** dataset, drop any non-matching cases from the left dataset

## Filtering joins

- `inner_join`: keep **only** data that matches in **both** datasets
- `full_join`: keep **all** the data in **both** datasets (also sometimes referred to as an *outer join*)

# Four types of joins

## Mutating joins

- `left_join`: keep all the data in the **left** dataset, drop any non-matching cases from the right dataset
- `right_join`: keep all the data in the **right** dataset, drop any non-matching cases from the left dataset

## Filtering joins

- `inner_join`: keep **only** data that matches in **both** datasets
- `full_join`: keep **all** the data in **both** datasets (also sometimes referred to as an *outer join*)

# Using joins to recode

Say you have a dataset like this

```
set.seed(1)
disab_codes <- c("00", "10", "20", "40", "43", "50", "60",
                 "70", "74", "80", "82", "90", "96", "98")
dis_tbl <- tibble(
  sid = 1:200,
  dis_code = sample(disab_codes, 200, replace = TRUE),
  score = as.integer(rnorm(200, 200, 10))
)
head(dis_tbl)
```

```
## # A tibble: 6 x 3
##   sid dis_code score
##   <int> <chr>   <int>
## 1     1 74      190
## 2     2 40      200
## 3     3 60      200
## 4     4 00      183
## 5     5 10      210
## 6     6 96      188
```



# Codes

Code	Disability
0	'Not Applicable'
10	'Intellectual Disability'
20	'Hearing Impairment'
40	'Visual Impairment'
43	'Deaf-Blindness'
50	'Communication Disorder'
60	'Emotional Disturbance'
70	'Orthopedic Impairment'
74	'Traumatic Brain Injury'
80	'Other Health Impairments'
82	'Autism Spectrum Disorder'
90	'Specific Learning Disability'
96	'Developmental Delay 0-2yr'
98	'Developmental Delay 3-4yr'

# Recode method

Using `case_when()`

```
dis_tbl %>%  
  mutate(disability = case_when(  
    dis_code == "10" ~ "Intellectual Disability",  
    dis_code == "20" ~ 'Hearing Impairment',  
    ...,  
    TRUE ~ "Not Applicable"  
  )  
)
```

# Join method

```
dis_code_tbl <- tibble(  
  dis_code = c(  
    "00", "10", "20", "40", "43", "50", "60",  
    "70", "74", "80", "82", "90", "96", "98"  
  ),  
  disability = c(  
    'Not Applicable', 'Intellectual Disability',  
    'Hearing Impairment', 'Visual Impairment',  
    'Deaf-Blindness', 'Communication Disorder',  
    'Emotional Disturbance', 'Orthopedic Impairment',  
    'Traumatic Brain Injury', 'Other Health Impairments',  
    'Autism Spectrum Disorder', 'Specific Learning Disability',  
    'Developmental Delay 0-2yr', 'Developmental Delay 3-4yr'  
  )  
)
```

## dis\_code\_tbl

```
## # A tibble: 14 x 2
##   dis_code disability
##   <chr>      <chr>
## 1 00        Not Applicable
## 2 10        Intellectual Disability
## 3 20        Hearing Impairment
## 4 40        Visual Impairment
## 5 43        Deaf-Blindness
## 6 50        Communication Disorder
## 7 60        Emotional Disturbance
## 8 70        Orthopedic Impairment
## 9 74        Traumatic Brain Injury
## 10 80       Other Health Impairments
## 11 82       Autism Spectrum Disorder
## 12 90       Specific Learning Disability
## 13 96       Developmental Delay 0-2yr
## 14 98       Developmental Delay 3-4yr
```

# Join the tables

```
left_join(dis_tbl, dis_code_tbl)
```

```
## Joining, by = "dis_code"
```

```
## # A tibble: 200 x 4
```

```
##       sid dis_code score disability
##   <int> <chr>    <int> <chr>
## 1     1  1 74      190 Traumatic Brain Injury
## 2     2  2 40      200 Visual Impairment
## 3     3  3 60      200 Emotional Disturbance
## 4     4  4 00      183 Not Applicable
## 5     5  5 10      210 Intellectual Disability
## 6     6  6 96      188 Developmental Delay 0-2yr
## 7     7  7 60      203 Emotional Disturbance
## 8     8  8 82      204 Autism Spectrum Disorder
## 9     9  9 98      201 Developmental Delay 3-4yr
## 10    10 10 10      198 Intellectual Disability
## # ... with 190 more rows
```

**Imperfect key match?**

# Consider the following

```
frl <- tibble(key = 1:3, frl = rbinom(3, 1, .5))  
sped <- tibble(key = c(1, 2, 4), sped = rbinom(3, 1, .5))
```

frl

```
## # A tibble: 3 x 2
```

```
##   key    frl
```

```
##   <int> <int>
```

```
## 1     1     0
```

```
## 2     2     1
```

```
## 3     3     0
```

sped

```
## # A tibble: 3 x 2
```

```
##   key    sped
```

```
##   <dbl> <int>
```

```
## 1     1     0
```

```
## 2     2     1
```

```
## 3     4     0
```

# Consider the following

## left\_join()?

```
left_join(frl, sped)
```


```
## # A tibble: 3 x 3
##   key   frl   sped
##   <dbl> <int> <int>
## 1     1     0     0
## 2     2     1     1
## 3     3     0    NA
```

## right\_join()?


```
right_join(frl, sped)
```

```
## # A tibble: 3 x 3
##   key   frl   sped
##   <dbl> <int> <int>
## 1     1     0     0
## 2     2     1     1
## 3     4    NA     0
```







Left



Right



Full




From [r4ds](#)


# Animations

**All of the following animations were created by Garrick Aden-Buie and can be found [here](#)**

`left_join(x, y)`




`right_join(x, y)`




# What if the key is not unique?

- Not an issue, as long as they are unique in one of the tables
  - In this case, it's called a one-to-many join



`left_join(x, y)`



# Example

## Student-level data

```
(stu <- tibble(  
  sid = 1:9,  
  scid = c(1, 1, 1, 1, 2, 2, 3, 3  
  score = c(10, 12, 15, 8, 9, 11
```

```
## # A tibble: 9 x 3  
##   sid  scid score  
##   <int> <dbl> <dbl>  
## 1     1     1    10  
## 2     2     1    12  
## 3     3     1    15  
## 4     4     1     8  
## 5     5     2     9  
## 6     6     2    11  
## 7     7     3    12  
## 8     8     3    15
```

## School-level data

```
(schl <- tibble(  
  scid = 1:3,  
  stu_tch_ratio = c(22.05, 31.14,  
  per_pupil_spending = c(15741.08  
  )  
)
```

```
## # A tibble: 3 x 3  
##   scid stu_tch_ratio per_pupil_spend  
##   <int>          <dbl>          <d  
## 1     1           22.0          157  
## 2     2           31.1          117  
## 3     3           24.9          130
```

# One to many

```
left_join(stu, schl)
```


```
## # A tibble: 9 x 5
##   sid  scid score stu_tch_ratio per_pupil_spending
##   <int> <dbl> <dbl>         <dbl>             <dbl>
## 1     1     1     10          22.0             15741.
## 2     2     1     12          22.0             15741.
## 3     3     1     15          22.0             15741.
## 4     4     1      8          22.0             15741.
## 5     5     2      9          31.1             11732.
## 6     6     2     11          31.1             11732.
## 7     7     3     12          24.9             13028.
## 8     8     3     15          24.9             13028.
## 9     9     3     17          24.9             13028.
```



# What if key is not unique to either table?

Generally this is an error

Result is probably not going to be what you want



# Example

```
seasonal_means <- tibble(  
  scid = rep(1:3, each = 3),  
  season = rep(c("fall", "winter", "spring"), 3),  
  mean = rnorm(3*3)  
)  
seasonal_means
```

```
## # A tibble: 9 x 3  
##   scid season  mean  
##   <int> <chr>   <dbl>  
## 1     1 fall    0.345  
## 2     1 winter  1.54  
## 3     1 spring -0.330  
## 4     2 fall    0.948  
## 5     2 winter -0.479  
## 6     2 spring -1.51  
## 7     3 fall    0.435  
## 8     3 winter -0.520  
## 9     3 spring -0.835
```

```
left_join(stu, seasonal_means)
```

```
## # A tibble: 27 x 5
##   sid  scid score season  mean
##   <int> <dbl> <dbl> <chr>   <dbl>
## 1     1     1     1    10 fall    0.345
## 2     1     1    10 winter  1.54
## 3     1     1    10 spring -0.330
## 4     2     1    12 fall    0.345
## 5     2     1    12 winter  1.54
## 6     2     1    12 spring -0.330
## 7     3     1    15 fall    0.345
## 8     3     1    15 winter  1.54
## 9     3     1    15 spring -0.330
## 10    4     1     8 fall    0.345
## # ... with 17 more rows
```

# How do we fix this?



In some cases, the solution is obvious, in others it is not

But **you must have at least one unique key** to join the datasets

# In this case

Move the dataset to wide before joining

## Move to wide

```
seasonal_means_wide <- seasonal_means %>%  
  pivot_wider(names_from = "season",  
              values_from = "mean")  
seasonal_means_wide
```

```
## # A tibble: 3 x 4  
##   scid  fall winter spring  
##   <int> <dbl>  <dbl>  <dbl>  
## 1     1  0.345  1.54   -0.330  
## 2     2  0.948 -0.479 -1.51  
## 3     3  0.435 -0.520 -0.835
```

We will cover this in Week 8

# Join

## One to many join

```
left_join(stu, seasonal_means_wide)
```

```
## # A tibble: 9 x 6
##   sid  scid score  fall winter spring
##   <int> <dbl> <dbl> <dbl>   <dbl>   <dbl>
## 1     1     1    10  0.345   1.54  -0.330
## 2     2     1    12  0.345   1.54  -0.330
## 3     3     1    15  0.345   1.54  -0.330
## 4     4     1     8  0.345   1.54  -0.330
## 5     5     2     9  0.948  -0.479 -1.51
## 6     6     2    11  0.948  -0.479 -1.51
## 7     7     3    12  0.435  -0.520 -0.835
## 8     8     3    15  0.435  -0.520 -0.835
## 9     9     3    17  0.435  -0.520 -0.835
```

# Another example

- Often you want to add summary info to your dataset
- You can do this easily with by piping arguments

## ECLS-K reminder

```
ecls
```

```
## # A tibble: 984 x 33
```

```
##   child_id teacher_id school_id k_type  school_type sex  ethnic famty
##   <chr>    <chr>      <chr>    <chr>    <chr>      <chr> <chr>  <chr>
## 1 0842021C 0842T02    0842    full-day public    male  BLACK 0~ BIOLO
## 2 0905002C 0905T01    0905    full-day private  male  ASIAN  BIOLO
## 3 0150012C 0150T01    0150    full-day private  female BLACK 0~ BIOLO
## 4 0556009C 0556T01    0556    full-day private  female HISPANI~ BIOLO
## 5 0089013C 0089T04    0089    full-day public    male  WHITE, ~ BIOLO
## 6 1217001C 1217T13    1217    half-day public    female NATIVE ~ BIOLO
## 7 1092008C 1092T01    1092    half-day public    female HISPANI~ BIOLO
## 8 0083007C 0083T16    0083    full-day public    male  WHITE, ~ BIOLO
## 9 1091005C 1091T02    1091    half-day private  male  WHITE, ~ BIOLO
## 10 2006006C 2006T01    2006    full-day private  male  WHITE, ~ BIOLO
```

# Compute group means

```
ec1s %>%  
  group_by(school_id) %>%  
  summarize(sch_pre_math = mean(T1MSCALE))
```

```
## # A tibble: 515 x 2  
##   school_id sch_pre_math  
##   <chr>      <dbl>  
## 1 0001      20.5  
## 2 0002      15.0  
## 3 0009      18.8  
## 4 0013      42.3  
## 5 0016      17.6  
## 6 0022      17.8  
## 7 0023      15.5  
## 8 0025      19.4  
## 9 0026      16.9  
## 10 0028      14.4  
## # ... with 505 more rows
```



# Join right within pipeline

```
ecls %>%  
  group_by(school_id) %>%  
  summarize(sch_pre_math = mean(T1MSCALE)) %>%  
  left_join(ecls) %>%  
  select(school_id:k_type) # Just for space
```

```
## # A tibble: 984 x 5
```

	school_id	sch_pre_math	child_id	teacher_id	k_type
	<chr>	<dbl>	<chr>	<chr>	<chr>
##	1 0001	20.5	0001010C	0001T01	full-day
##	2 0002	15.0	0002010C	0002T01	half-day
##	3 0009	18.8	0009026C	0009T01	half-day
##	4 0009	18.8	0009014C	0009T02	half-day
##	5 0009	18.8	0009005C	0009T01	half-day
##	6 0013	42.3	0013003C	0013T01	full-day
##	7 0016	17.6	0016004C	0016T01	half-day
##	8 0016	17.6	0016009C	0016T01	half-day
##	9 0022	17.8	0022005C	0022T01	half-day
##	10 0022	17.8	0022014C	0022T03	half-day

```
## # ... with 974 more rows
```

# Default join behavior

By default, the `_join` functions will use all columns with common names as keys

```
flights2 <- flights %>%  
  select(year:day, hour, origin, dest, tailnum, carrier)  
flights2[1:2, ]
```

```
## # A tibble: 2 x 8
```

	year	month	day	hour	origin	dest	tailnum	carrier
	<int>	<int>	<int>	<dbl>	<chr>	<chr>	<chr>	<chr>
1	2013	1	1	5	EWB	IAH	N14228	UA
2	2013	1	1	5	LGA	IAH	N24211	UA

```
weather[1:2, ]
```

```
## # A tibble: 2 x 15
```

	origin	year	month	day	hour	temp	dewp	humid	wind_dir	wind_speed	wind
	<chr>	<int>	<int>	<int>	<int>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	
1	EWB	2013	1	1	1	39.0	26.1	59.4	270	10.4	
2	EWB	2013	1	1	2	39.0	27.0	61.6	250	8.06	

## # ... with 1 more variable: time\_hour <dtm>

```
left_join(flights2, weather)
```

```
## Joining, by = c("year", "month", "day", "hour", "origin")
```

```
## # A tibble: 336,776 x 18
```

```
##   year month   day hour origin dest tailnum carrier  temp  dewp humid w
##   <int> <int> <int> <dbl> <chr> <chr> <chr>   <chr>   <dbl> <dbl> <dbl> w
## 1  2013     1     1     5 EWR   IAH   N14228 UA      39.0  28.0  64.4
## 2  2013     1     1     5 LGA   IAH   N24211 UA      39.9  25.0  54.8
## 3  2013     1     1     5 JFK   MIA   N619AA AA      39.0  27.0  61.6
## 4  2013     1     1     5 JFK   BQN   N804JB B6      39.0  27.0  61.6
## 5  2013     1     1     6 LGA   ATL   N668DN DL      39.9  25.0  54.8
## 6  2013     1     1     5 EWR   ORD   N39463 UA      39.0  28.0  64.4
## 7  2013     1     1     6 EWR   FLL   N516JB B6      37.9  28.0  67.2
## 8  2013     1     1     6 LGA   IAD   N829AS EV      39.9  25.0  54.8
## 9  2013     1     1     6 JFK   MCO   N593JB B6      37.9  27.0  64.3
## 10 2013     1     1     6 LGA   ORD   N3ALAA AA      39.9  25.0  54.8
## # ... with 336,766 more rows, and 4 more variables: precip <dbl>, pressure
## #   time_hour <dtm>
```

# Use only certain keys?

If we were joining *flights2* and *planes*, we would not want to use the `year` variable in the join, because **it means different things in each dataset**

```
head(planes)
```

```
## # A tibble: 6 x 9
```

	tailnum	year	type	manufacturer	model	engines
	<chr>	<int>	<chr>	<chr>	<chr>	<int>
## 1	N10156	2004	Fixed wing multi engine	EMBRAER	EMB-145XR	2
## 2	N102UW	1998	Fixed wing multi engine	AIRBUS INDUSTRIE	A320-214	2
## 3	N103US	1999	Fixed wing multi engine	AIRBUS INDUSTRIE	A320-214	2
## 4	N104UW	1999	Fixed wing multi engine	AIRBUS INDUSTRIE	A320-214	2
## 5	N10575	2002	Fixed wing multi engine	EMBRAER	EMB-145LR	2
## 6	N105UW	1999	Fixed wing multi engine	AIRBUS INDUSTRIE	A320-214	2

# Specify `*_join()` keys

Specify the variables with `by`

```
left_join(flights2, planes, by = "tailnum")
```

```
## # A tibble: 336,776 x 16
```

```
##   year.x month   day   hour origin dest tailnum carrier year.y type  manu
##   <int> <int> <int> <dbl> <chr> <chr> <chr>   <chr>   <int> <chr> <chr>
## 1  2013     1     1     5   EWR   IAH   N14228   UA      1999 Fixed~ BOEI
## 2  2013     1     1     5   LGA   IAH   N24211   UA      1998 Fixed~ BOEI
## 3  2013     1     1     5   JFK   MIA   N619AA   AA      1990 Fixed~ BOEI
## 4  2013     1     1     5   JFK   BQN   N804JB   B6      2012 Fixed~ AIRB
## 5  2013     1     1     6   LGA   ATL   N668DN   DL      1991 Fixed~ BOEI
## 6  2013     1     1     5   EWR   ORD   N39463   UA      2012 Fixed~ BOEI
## 7  2013     1     1     6   EWR   FLL   N516JB   B6      2000 Fixed~ AIRB
## 8  2013     1     1     6   LGA   IAD   N829AS   EV      1998 Fixed~ CANA
## 9  2013     1     1     6   JFK   MCO   N593JB   B6      2004 Fixed~ AIRB
## 10 2013     1     1     6   LGA   ORD   N3ALAA   AA       NA <NA>  <NA>
## # ... with 336,766 more rows, and 1 more variable: engine <chr>
```

# Specify `*_join()` keys

I like to *always* specify the `by` vars

Makes intent explicit

Helps me review my own code

# Mismatched key names

What if you had data to merge like this?

```
names(schl)[1] <- "school_id"
schl
```

```
## # A tibble: 3 x 3
##   school_id stu_tch_ratio per_pupil_spending
##   <int>      <dbl>      <dbl>
## 1      1      22.0      15741.
## 2      2      31.1      11732.
## 3      3      24.9      13028.
```

```
stu
```

```
## # A tibble: 9 x 3
##   sid  scid score
##   <int> <dbl> <dbl>
## 1      1      1    10
## 2      2      1    12
## 3      3      1    15
## 4      4      1     8
```

# Join with mismatched key names

```
left_join(stu, sch1, by = c("scid" = "school_id"))
```

```
## # A tibble: 9 x 5
##   sid  scid score stu_tch_ratio per_pupil_spending
##   <int> <dbl> <dbl>         <dbl>             <dbl>
## 1     1     1     10          22.0             15741.
## 2     2     1     12          22.0             15741.
## 3     3     1     15          22.0             15741.
## 4     4     1      8          22.0             15741.
## 5     5     2      9          31.1             11732.
## 6     6     2     11          31.1             11732.
## 7     7     3     12          24.9             13028.
## 8     8     3     15          24.9             13028.
## 9     9     3     17          24.9             13028.
```



**Next time**

# Before next class

- Homework
  - **Homework 5**
- Reading
  - [R4DS 29](#)
- Complete
  - [Markdown Tutorial](#)

# Homework 5



